

```

1
2 ## install required packages
3 !pip install swig
4 !pip install wrds
5 !pip install pyportfolioopt
6 ## install finrl library
7 !pip install -q condacolab
8 import condacolab
9 condacolab.install()
10 !apt-get update -y -qq && apt-get install -y -qq cmake libopenmpi-dev python3-dev zlib1g-dev libgl1-mesa-glx swig
11 !pip install git+https://github.com/AI4Finance-Foundation/FinRL.git

→ Requirement already satisfied: swig in /usr/local/lib/python3.10/site-packages (4.2.1)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. ] Requirement already satisfied: wrds in /usr/local/lib/python3.10/site-packages (3.2.0)
Requirement already satisfied: numpy<1.27,>=1.26 in /usr/local/lib/python3.10/site-packages (from wrds) (1.26.4)
Requirement already satisfied: packaging<23.3 in /usr/local/lib/python3.10/site-packages (from wrds) (23.2)
Requirement already satisfied: pandas<2.3,>=2.2 in /usr/local/lib/python3.10/site-packages (from wrds) (2.2.2)
Requirement already satisfied: psycopg2-binary<2.10,>=2.9 in /usr/local/lib/python3.10/site-packages (from wrds) (2.9.9)
Requirement already satisfied: scipy<1.13,>=1.12 in /usr/local/lib/python3.10/site-packages (from wrds) (1.12.0)
Requirement already satisfied: sqlalchemy<2.1,>=2 in /usr/local/lib/python3.10/site-packages (from wrds) (2.0.31)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/site-packages (from pandas<2.3,>=2.2->wrds) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/site-packages (from pandas<2.3,>=2.2->wrds) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/site-packages (from pandas<2.3,>=2.2->wrds) (2024.1)
Requirement already satisfied: typing-extensions>=4.6.0 in /usr/local/lib/python3.10/site-packages (from sqlalchemy<2.1,>=2->wrds) (4.6.1)
Requirement already satisfied: greenlet!=0.4.17 in /usr/local/lib/python3.10/site-packages (from sqlalchemy<2.1,>=2->wrds) (3.0.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas<2.3,>=2.2->wrds) (1.5.2)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. ] Requirement already satisfied: pyportfolioopt in /usr/local/lib/python3.10/site-packages (1.5.5)
Requirement already satisfied: cvxpy<2.0.0,>=1.1.19 in /usr/local/lib/python3.10/site-packages (from pyportfolioopt) (1.5.2)
Requirement already satisfied: numpy<2.0.0,>=1.22.4 in /usr/local/lib/python3.10/site-packages (from pyportfolioopt) (1.26.4)
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.10/site-packages (from pyportfolioopt) (2.2.2)
Requirement already satisfied: scipy<2.0,>=1.3 in /usr/local/lib/python3.10/site-packages (from pyportfolioopt) (1.12.0)
Requirement already satisfied: osqp>=0.6.2 in /usr/local/lib/python3.10/site-packages (from cvxpy<2.0.0,>=1.1.19->pyportfolioopt) (0.6.2)
Requirement already satisfied: ecos>=2 in /usr/local/lib/python3.10/site-packages (from cvxpy<2.0.0,>=1.1.19->pyportfolioopt) (2.0.12)
Requirement already satisfied: clarabel>=0.5.0 in /usr/local/lib/python3.10/site-packages (from cvxpy<2.0.0,>=1.1.19->pyportfolioopt) (0.5.0)
Requirement already satisfied: scs>=3.2.4.post1 in /usr/local/lib/python3.10/site-packages (from cvxpy<2.0.0,>=1.1.19->pyportfolioopt) (3.2.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/site-packages (from pandas>=0.19->pyportfolioopt) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/site-packages (from pandas>=0.19->pyportfolioopt) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/site-packages (from pandas>=0.19->pyportfolioopt) (2024.1)
Requirement already satisfied: qldl in /usr/local/lib/python3.10/site-packages (from osqp>=0.6.2->cvxpy<2.0.0,>=1.1.19->pyportfolioopt) (0.6.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/site-packages (from python-dateutil>=2.8.2->pandas>=0.19->pyportfolioopt) (1.5.2)
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. ] WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. ] ✨ Everything looks OK!
Collecting git+https://github.com/AI4Finance-Foundation/FinRL.git
  Cloning https://github.com/AI4Finance-Foundation/FinRL.git to /tmp/pip-req-build-afqd_zji
    Running command git clone --filter=blob:none --quiet https://github.com/AI4Finance-Foundation/FinRL.git /tmp/pip-req-build-afqd_zji
  Resolved https://github.com/AI4Finance-Foundation/FinRL.git to commit c7c0429a14f94fd2a460558bb49240e184f8d32d
  Installing build dependencies ... done
  Getting requirements to build wheel ... done
  Preparing metadata (pyproject.toml) ... done
Collecting elegantrl@ git+https://github.com/AI4Finance-Foundation/ElegantRL.git#egg=elegantrl (from finrl==0.3.6)
  Cloning https://github.com/AI4Finance-Foundation/ElegantRL.git to /tmp/pip-install-0325irt3/elegantrl_666fe7eb368b4adebf82c8dd4668c
  Running command git clone --filter=blob:none --quiet https://github.com/AI4Finance-Foundation/ElegantRL.git /tmp/pip-install-0325irt3/elegantrl_666fe7eb368b4adebf82c8dd4668c
  Resolved https://github.com/AI4Finance-Foundation/ElegantRL.git to commit 6ead2cbd17e4ecd3942fc8a394719db1515cf13f
  Preparing metadata (setup.py) ... done
Requirement already satisfied: alpaca-trade-api<4,>=3 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (3.2.0)
Requirement already satisfied: ccxt<4,>=3 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (3.1.60)
Requirement already satisfied: exchange-calendars<5,>=4 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (4.5.5)
Requirement already satisfied: jqdatasdk<2,>=1 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (1.9.5)
Requirement already satisfied: pyfolio<0.10,>=0.9 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (0.9.2)
Requirement already satisfied: pyportfolioopt<2,>=1 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (1.5.5)
Requirement already satisfied: ray<3,>=2 in /usr/local/lib/python3.10/site-packages (from ray[default,tune]<3,>=2->finrl==0.3.6) (2.2.0)
Requirement already satisfied: scikit-learn<2,>=1 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (1.5.1)
Requirement already satisfied: stable-baselines3>=2.0.0a5 in /usr/local/lib/python3.10/site-packages (from stable-baselines3[extra]>=2.0.0a5) (2.0.0a5)
Requirement already satisfied: stockstats<0.6,>=0.5 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (0.5.4)
Requirement already satisfied: wrds<4,>=3 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (3.2.0)
Requirement already satisfied: vfinance<0.3,>=0.2 in /usr/local/lib/python3.10/site-packages (from finrl==0.3.6) (0.2.40)

```

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib
4 import matplotlib.pyplot as plt
5 # matplotlib.use('Agg')
6 import datetime
7
8 %matplotlib inline
9 from finrl.meta.preprocessor.yahoodownloader import YahooDownloader
10 from finrl.meta.preprocessors import FeatureEngineer, data_split
11 from finrl.meta.env_stock_trading.env_stocktrading import StockTradingEnv
12 from finrl.agents.stablebaselines3.models import DRLAgent
13 from stable_baselines3.common.logger import configure
14 from finrl.meta.data_processor import DataProcessor
15
16 from finrl.plot import backtest_stats, backtest_plot, get_daily_return, get_baseline
17 from pprint import pprint
18
19 import sys
20 sys.path.append("../FinRL")
21
22 import iterools

→ /usr/local/lib/python3.10/site-packages/pyfolio/pos.py:26: UserWarning: Module "zipline.assets" not found; multipliers will not be applied
  warnings.warn(

```



```

1 from finrl import config
2 from finrl import config_tickers
3 import os
4 from finrl.main import check_and_make_directories
5 from finrl.config import (
6     DATA_SAVE_DIR,
7     TRAINED_MODEL_DIR,
8     TENSORBOARD_LOG_DIR,
9     RESULTS_DIR,
10    INDICATORS,
11    TRAIN_START_DATE,
12    TRAIN_END_DATE,
13    TEST_START_DATE,
14    TEST_END_DATE,
15    TRADE_START_DATE,
16    TRADE_END_DATE,
17 )
18 check_and_make_directories([DATA_SAVE_DIR, TRAINED_MODEL_DIR, TENSORBOARD_LOG_DIR, RESULTS_DIR])

```

```

1 # from config.py, TRAIN_START_DATE is a string
2 TRAIN_START_DATE
3 # from config.py, TRAIN_END_DATE is a string
4 TRAIN_END_DATE

```

→ '2020-07-31'

```

1 TRAIN_START_DATE = '2010-01-01'
2 TRAIN_END_DATE = '2021-10-01'
3 TRADE_START_DATE = '2021-10-01'
4 TRADE_END_DATE = '2023-03-01'

```

```

1 df = YahooDownloader(start_date = TRAIN_START_DATE,
2                      end_date = TRADE_END_DATE,
3                      ticker_list = config_tickers.DOW_30_TICKER).fetch_data()

```

```

→ [*****100%*****] 1 of 1 completed

```

```
[*****100%*****] 1 of 1 completed
Shape of DataFrame: (97013, 8)
```

```
1 print(config_tickers.DOW_30_TICKER)
```

```
→ ['AXP', 'AMGN', 'AAPL', 'BA', 'CAT', 'CSCO', 'CVX', 'GS', 'HD', 'HON', 'IBM', 'INTC', 'JNJ', 'KO', 'JPM', 'MCD', 'MMM', 'MRK', 'MSFT', 'PFE', 'TGT', 'VZ', 'WMT', 'XOM']
```

```
1 df.shape
```

```
→ (97013, 8)
```

```
1 df.sort_values(['date','tic'],ignore_index=True).head()
```

	date	open	high	low	close	volume	tic	day
0	2010-01-04	7.622500	7.660714	7.585000	6.461976	493729600	AAPL	0
1	2010-01-04	56.630001	57.869999	56.560001	41.200794	5277400	AMGN	0
2	2010-01-04	40.810001	41.099998	40.389999	33.090431	6894300	AXP	0
3	2010-01-04	55.720001	56.389999	54.799999	43.777538	6186700	BA	0
4	2010-01-04	57.650002	59.189999	57.509998	40.190208	7325600	CAT	0

```
1 fe = FeatureEngineer(
2     use_technical_indicator=True,
3     tech_indicator_list = INDICATORS,
4     use_vix=True,
5     use_turbulence=True,
6     user_defined_feature = False)
7
8 processed = fe.preprocess_data(df)
```

```
→ Successfully added technical indicators
```

```
[*****100%*****] 1 of 1 completed
```

```
Shape of DataFrame: (3310, 8)
```

```
Successfully added vix
```

```
Successfully added turbulence index
```

```
1 list_ticker = processed["tic"].unique().tolist()
2 list_date = list(pd.date_range(processed['date'].min(),processed['date'].max()).astype(str))
3 combination = list(itertools.product(list_date,list_ticker))
4
5 processed_full = pd.DataFrame(combination,columns=["date","tic"]).merge(processed,on=["date","tic"],how="left")
6 processed_full = processed_full[processed_full['date'].isin(processed['date'])]
7 processed_full = processed_full.sort_values(['date','tic'])
8
9 processed_full = processed_full.fillna(0)
```

```
1 processed_full.sort_values(['date','tic'],ignore_index=True).head(10)
```

	date	tic	open	high	low	close	volume	day	macd	boll_ub	boll_lb	rsi_30	cci_30	dx_30	close_
0	2010-01-04	AAPL	7.622500	7.660714	7.585000	6.461976	493729600.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	6
1	2010-01-04	AMGN	56.630001	57.869999	56.560001	41.200794	5277400.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	41
2	2010-01-04	AXP	40.810001	41.099998	40.389999	33.090431	6894300.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	33
3	2010-01-04	BA	55.720001	56.389999	54.799999	43.777538	6186700.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	43
4	2010-01-04	CAT	57.650002	59.189999	57.509998	40.190208	7325600.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	40
5	2010-01-04	CRM	18.652500	18.882500	18.547501	18.680597	7906000.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	18
6	2010-01-04	CSCO	24.110001	24.840000	24.010000	16.743504	59853700.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	16
7	2010-01-04	CVX	78.199997	79.199997	78.160004	44.538490	10173800.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	44
8	2010-01-04	DIS	32.500000	32.750000	31.870001	27.843174	13700400.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	27
9	2010-01-04	GS	170.050003	174.250000	169.509995	134.760635	9135000.0	0.0	0.0	6.483364	6.451761	100.0	66.666667	100.0	134

```
1 mvo_df = processed_full.sort_values(['date','tic'],ignore_index=True)[['date','tic','close']]
```

```
1 train = data_split(processed_full, TRAIN_START_DATE,TRAIN_END_DATE)
2 trade = data_split(processed_full, TRADE_START_DATE,TRADE_END_DATE)
3 print(len(train))
4 print(len(trade))
```

85753  
10237

```
1 train.tail()
```

	date	tic	open	high	low	close	volume	day	macd	boll_ub	boll_lb	rsi_30	cci_30
2956	2021-09-30	UNH	401.489990	403.489990	390.459991	375.759766	3779900.0	3.0	-4.265029	411.094077	379.372033	40.895406	-222.754203
2956	2021-09-30	V	227.580002	228.789993	222.630005	218.139252	7128500.0	3.0	-1.523925	226.440142	214.446931	44.078990	-54.509447
2956	2021-09-30	VZ	54.500000	54.509998	54.000000	45.622719	18736600.0	3.0	-0.217510	46.664909	45.373700	41.824968	-102.367070
2956	2021-09-30	WBA	48.790001	48.930000	46.919998	40.519119	6449400.0	3.0	-0.233708	44.735839	40.187235	44.613740	-106.638112
2956	2021-09-30	WMT	46.880001	47.243332	46.416668	44.566532	22457700.0	3.0	-0.509083	47.853937	44.408838	40.165770	-151.493845

```
1 trade.head()
```

	date	tic	open	high	low	close	volume	day	macd	boll_ub	boll_lb	rsi_30	cci_30
0	2021-10-01	AAPL	141.899994	142.919998	139.110001	140.462982	94639600.0	4.0	-1.692147	154.348516	136.219347	46.927719	-142.229561 4
0	2021-10-01	AMGN	213.589996	214.610001	210.800003	195.822235	2629400.0	4.0	-2.975414	204.393145	191.531747	40.408523	-97.013199 3
0	2021-10-01	AXP	168.500000	175.119995	168.479996	167.727448	3956000.0	4.0	2.242077	171.823829	147.181359	56.265110	117.437341 1
0	2021-10-01	BA	222.850006	226.720001	220.600006	226.000000	9113600.0	4.0	0.730320	226.909442	205.727561	51.614047	116.649440
0	2021-10-01	CAT	192.899994	195.869995	191.240005	183.514008	3695500.0	4.0	-3.554818	200.903604	177.171314	41.999445	-112.169824 3

## 1 INDICATORS

```
→ ['macd',
 'boll_ub',
 'boll_lb',
 'rsi_30',
 'cci_30',
 'dx_30',
 'close_30_sma',
 'close_60_sma']
```

```
1 stock_dimension = len(train.tic.unique())
2 state_space = 1 + 2*stock_dimension + len(INDICATORS)*stock_dimension
3 print(f"Stock Dimension: {stock_dimension}, State Space: {state_space}")
4
```

→ Stock Dimension: 29, State Space: 291

```
1 buy_cost_list = sell_cost_list = [0.001] * stock_dimension
2 num_stock_shares = [0] * stock_dimension
3
4 env_kwargs = {
5     "hmax": 100,
6     "initial_amount": 1000000,
7     "num_stock_shares": num_stock_shares,
8     "buy_cost_pct": buy_cost_list,
9     "sell_cost_pct": sell_cost_list,
10    "state_space": state_space,
11    "stock_dim": stock_dimension,
12    "tech_indicator_list": INDICATORS,
13    "action_space": stock_dimension,
14    "reward_scaling": 1e-4
15 }
16
17
18 e_train_gym = StockTradingEnv(df = train, **env_kwargs)
```

```
1 env_train, _ = e_train_gym.get_sb_env()
2 print(type(env_train))
```

→ <class 'stable\_baselines3.common.vec\_env.dummy\_vec\_env.DummyVecEnv'>

```
1 agent = DRLAgent(env = env_train)
2
3 if_using_a2c = True
4 if_using_ddpg = True
5 if_using_ppo = True
6 if_using_td3 = True
7 if_using_sac = True
```

```

1 agent = DRLAgent(env = env_train)
2 model_a2c = agent.get_model("a2c")
3
4 if if_using_a2c:
5     # set up logger
6     tmp_path = RESULTS_DIR + '/a2c'
7     new_logger_a2c = configure(tmp_path, ["stdout", "csv", "tensorboard"])
8     # Set new logger
9     model_a2c.set_logger(new_logger_a2c)

```

↳ { 'n\_steps': 5, 'ent\_coef': 0.01, 'learning\_rate': 0.0007}  
Using cpu device  
Logging to results/a2c

```

1 trained_a2c = agent.train_model(model=model_a2c,
2                                 tb_log_name='a2c',
3                                 total_timesteps=50000) if if_using_a2c else None

```

time/	
fps	72
iterations	100
time_elapsed	6
total_timesteps	500
train/	
entropy_loss	-41.2
explained_variance	0.0459
learning_rate	0.0007
n_updates	99
policy_loss	-75.3
reward	-0.43437988
std	1
value_loss	4.29

  

time/	
fps	71
iterations	200
time_elapsed	14
total_timesteps	1000
train/	
entropy_loss	-41.3
explained_variance	0
learning_rate	0.0007
n_updates	199
policy_loss	-84.3
reward	0.8485453
std	1
value_loss	8.62

  

time/	
fps	70
iterations	300
time_elapsed	21
total_timesteps	1500
train/	
entropy_loss	-41.3
explained_variance	0.053
learning_rate	0.0007
n_updates	299
policy_loss	-50.4
reward	-1.8449868
std	1
value_loss	4.46

  

time/	
fps	70
iterations	400
time_elapsed	28
total_timesteps	2000
train/	
entropy_loss	-41.3
explained_variance	5.96e-08
learning_rate	0.0007

## ▼ Agent 2: DDPG

```

1 agent = DRLAgent(env = env_train)
2 model_ddpg = agent.get_model("ddpg")
3
4 if if_using_ddpg:
5     # set up logger
6     tmp_path = RESULTS_DIR + '/ddpg'
7     new_logger_ddpg = configure(tmp_path, ["stdout", "csv", "tensorboard"])
8     # Set new logger
9     model_ddpg.set_logger(new_logger_ddpg)

{'batch_size': 128, 'buffer_size': 50000, 'learning_rate': 0.001}
Using cpu device
Logging to results/ddpg

1 trained_ddpg = agent.train_model(model=model_ddpg,
2                                     tb_log_name='ddpg',
3                                     total_timesteps=50000) if if_using_ddpg else None

day: 2956, episode: 20
begin_total_asset: 1000000.00
end_total_asset: 4798598.58
total_reward: 3798598.58
total_cost: 1093.97
total_trades: 50253
Sharpe: 0.873
=====


|                 |           |
|-----------------|-----------|
| time/           |           |
| episodes        | 4         |
| fps             | 20        |
| time_elapsed    | 573       |
| total_timesteps | 11828     |
| train/          |           |
| actor_loss      | -82.1     |
| critic_loss     | 19.9      |
| learning_rate   | 0.001     |
| n_updates       | 11727     |
| reward          | -7.658464 |


=====



|                 |           |
|-----------------|-----------|
| time/           |           |
| episodes        | 8         |
| fps             | 20        |
| time_elapsed    | 1157      |
| total_timesteps | 23656     |
| train/          |           |
| actor_loss      | -49.5     |
| critic_loss     | 2.51      |
| learning_rate   | 0.001     |
| n_updates       | 23555     |
| reward          | -7.658464 |


=====

day: 2956, episode: 30
begin_total_asset: 1000000.00
end_total_asset: 4837059.37
total_reward: 3837059.37
total_cost: 999.00
total_trades: 50214
Sharpe: 0.864
=====


|                 |           |
|-----------------|-----------|
| time/           |           |
| episodes        | 12        |
| fps             | 20        |
| time_elapsed    | 1754      |
| total_timesteps | 35484     |
| train/          |           |
| actor_loss      | -35       |
| critic_loss     | 2.61      |
| learning_rate   | 0.001     |
| n_updates       | 35383     |
| reward          | -7.658464 |


=====



|          |    |
|----------|----|
| time/    |    |
| episodes | 16 |


```

## Agent 3:PPO

```

1 agent = DRLAgent(env = env_train)
2 PPO_PARAMS = {
3     "n_steps": 2048,
4     "ent_coef": 0.01,
5     "learning_rate": 0.00025,
6     "batch_size": 128,
7 }
8 model_ppo = agent.get_model("ppo",model_kwargs = PPO_PARAMS)
9
10 if if_using_ppo:
11     # set up logger
12     tmp_path = RESULTS_DIR + '/ppo'
13     new_logger_ppo = configure(tmp_path, ["stdout", "csv", "tensorboard"])
14     # Set new logger
15     model_ppo.set_logger(new_logger_ppo)

↳ {'n_steps': 2048, 'ent_coef': 0.01, 'learning_rate': 0.00025, 'batch_size': 128}
Using cpu device
Logging to results/ppo

```

```

1 trained_ppo = agent.train_model(model=model_ppo,
2                                 tb_log_name='ppo',
3                                 total_timesteps=50000) if if_using_ppo else None

```

Time /		
fps	74	
iterations	1	
time_elapsed	27	
total_timesteps	2048	
Train /		
reward	2.0119593	

  

Time /		
fps	72	
iterations	2	
time_elapsed	56	
total_timesteps	4096	
Train /		
approx_kl	0.020232465	
clip_fraction	0.237	
clip_range	0.2	
entropy_loss	-41.2	
explained_variance	-0.0223	
learning_rate	0.00025	
loss	5.72	
n_updates	10	
policy_gradient_loss	-0.0306	
reward	-0.62762034	
std	1	
value_loss	13.3	

  

Time /		
fps	72	
iterations	3	
time_elapsed	85	
total_timesteps	6144	
Train /		
approx_kl	0.014667155	
clip_fraction	0.199	
clip_range	0.2	
entropy_loss	-41.3	
explained_variance	-0.00104	
learning_rate	0.00025	
loss	21.5	
n_updates	20	
policy_gradient_loss	-0.0155	
reward	2.111274	
std	1	
value_loss	40.8	

  

Time /		
fps	71	
iterations	4	
time_elapsed	113	
total_timesteps	8192	
Train /		
approx_kl	0.018692512	
clip_fraction	0.178	

## Agent 4: TD3

```

1 agent = DRLAgent(env = env_train)
2 TD3_PARAMS = {"batch_size": 100,
3                 "buffer_size": 1000000,
4                 "learning_rate": 0.001}
5
6 model_td3 = agent.get_model("td3",model_kwarg = TD3_PARAMS)
7
8 if if_using_td3:
9     # set up logger
10    tmp_path = RESULTS_DIR + '/td3'
11    new_logger_td3 = configure(tmp_path, ["stdout", "csv", "tensorboard"])
12    # Set new logger
13    model_td3.set_logger(new_logger_td3)

```

→ {'batch\_size': 100, 'buffer\_size': 1000000, 'learning\_rate': 0.001}  
 Using cpu device  
 Logging to results/td3

```

1 trained_td3 = agent.train_model(model=model_td3,
2                                 tb_log_name='td3',
3                                 total_timesteps=50000) if if_using_td3 else None

```

time/	
episodes	4
fps	21
time_elapsed	553
total_timesteps	11828
train/	
actor_loss	7.73
critic_loss	515
learning_rate	0.001
n_updates	11727
reward	-6.0941057

---

day: 2956, episode: 60
begin_total_asset: 1000000.00
end_total_asset: 3637091.88
total_reward: 2637091.88
total_cost: 999.00
total_trades: 41365
Sharpe: 0.699

---

time/	
episodes	8
fps	20
time_elapsed	1134
total_timesteps	23656
train/	
actor_loss	3.36
critic_loss	16.2
learning_rate	0.001
n_updates	23555
reward	-6.0941057

---

time/	
episodes	12
fps	20
time_elapsed	1715
total_timesteps	35484
train/	
actor_loss	4.66
critic_loss	6.6
learning_rate	0.001
n_updates	35383
reward	-6.0941057

---

time/	
episodes	16
fps	20
time_elapsed	2302
total_timesteps	47312
train/	

actor_loss	13.9
critic_loss	12.3
learning_rate	0.001
n_updates	47211

## Agent 5: SAC

```

1 agent = DRLAgent(env = env_train)
2 SAC_PARAMS = {
3     "batch_size": 128,
4     "buffer_size": 100000,
5     "learning_rate": 0.0001,
6     "learning_starts": 100,
7     "ent_coef": "auto_0.1",
8 }
9
10 model_sac = agent.get_model("sac",model_kwargs = SAC_PARAMS)
11
12 if if_using_sac:
13     # set up logger
14     tmp_path = RESULTS_DIR + '/sac'
15     new_logger_sac = configure(tmp_path, ["stdout", "csv", "tensorboard"])
16     # Set new logger
17     model_sac.set_logger(new_logger_sac)

{'batch_size': 128, 'buffer_size': 100000, 'learning_rate': 0.0001, 'learning_starts': 100, 'ent_coef': 'auto_0.1'}
Using cpu device
Logging to results/sac

```

```

1 trained_sac = agent.train_model(model=model_sac,
2                                 tb_log_name='sac',
3                                 total_timesteps=50000) if if_using_sac else None

```

time/	
episodes	4
fps	20
time_elapsed	569
total_timesteps	11828
train/	
actor_loss	828
critic_loss	2.33e+03
ent_coef	0.137
ent_coef_loss	-88.5
learning_rate	0.0001
n_updates	11727
reward	-9.536868

  

time/	
episodes	8
fps	20
time_elapsed	1155
total_timesteps	23656
train/	
actor_loss	334
critic_loss	76.5
ent_coef	0.0434
ent_coef_loss	-96.3
learning_rate	0.0001
n_updates	23555
reward	-9.587751

  

day: 2956, episode: 80	
begin_total_asset:	1000000.00
end_total_asset:	5420950.10
total_reward:	4420950.10
total_cost:	4211.44
total_trades:	46297
Sharpe:	0.764

  

time/	
episodes	12
fps	20
time_elapsed	1740
total_timesteps	35484
train/	

actor_loss	170
critic_loss	31.9
ent_coef	0.014
ent_coef_loss	-76.1
learning_rate	0.0001
n_updates	35383
reward	-8.331024

  

time/	
episodes	16
fps	20

## ▼ In Sample Performance

Assume that the initial capital is \$1,000,000.

### Set turbulence threshold

Set the turbulence threshold to be greater than the maximum of insample turbulence data. If current turbulence index is greater than the threshold, then we assume that the current market is volatile

```
1 data_risk_indicator = processed_full[(processed_full.date<TRAIN_END_DATE) & (processed_full.date>=TRAIN_START_DATE)]
2 insample_risk_indicator = data_risk_indicator.drop_duplicates(subset=['date'])
```

```
1
2 insample_risk_indicator.vix.describe()
```

```
→ count    2957.000000
mean      18.105293
std       7.272476
min       9.140000
25%      13.370000
50%      16.209999
75%      20.629999
max      82.690002
Name: vix, dtype: float64
```

```
1 insample_risk_indicator.vix.quantile(0.996)
```

```
→ 57.212001831054636
```

```
1 insample_risk_indicator.turbulence.describe()
```

```
→ count    2957.000000
mean      34.139588
std       43.879186
min       0.000000
25%      14.613418
50%      23.644978
75%      38.292511
max      652.504962
Name: turbulence, dtype: float64
```

```
1
2 insample_risk_indicator.turbulence.quantile(0.996)
```

```
→ 291.7246732286184
```

## ▼ Trading (Out-of-sample Performance)

We update periodically in order to take full advantage of the data, e.g., retrain quarterly, monthly or weekly. We also tune the parameters along the way, in this notebook we use the in-sample data from 2009-01 to 2020-07 to tune the parameters once, so there is some alpha decay here as the length of trade date extends.

Numerous hyperparameters – e.g. the learning rate, the total number of samples to train on – influence the learning process and are usually determined by testing some variations.

```

1
2 e_trade_gym = StockTradingEnv(df = trade, turbulence_threshold = 70,risk_indicator_col='vix', **env_kwargs)
3 # env_trade, obs_trade = e_trade_gym.get_sb_env()

```

```
1 trade.head()
```

	date	tic	open	high	low	close	volume	day	macd	boll_ub	boll_lb	rsi_30	cci_30
0	2021-10-01	AAPL	141.899994	142.919998	139.110001	140.462982	94639600.0	4.0	-1.692147	154.348516	136.219347	46.927719	-142.229561 4
0	2021-10-01	AMGN	213.589996	214.610001	210.800003	195.822235	2629400.0	4.0	-2.975414	204.393145	191.531747	40.408523	-97.013199 3
0	2021-10-01	AXP	168.500000	175.119995	168.479996	167.727448	3956000.0	4.0	2.242077	171.823829	147.181359	56.265110	117.437341 1
0	2021-10-01	BA	222.850006	226.720001	220.600006	226.000000	9113600.0	4.0	0.730320	226.909442	205.727561	51.614047	116.649440
0	2021-10-01	CAT	192.899994	195.869995	191.240005	183.514008	3695500.0	4.0	-3.554818	200.903604	177.171314	41.999445	-112.169824 3

```

1
2 trained_moedl = trained_a2c
3 df_account_value_a2c, df_actions_a2c = DRLAgent.DRL_prediction(
4     model=trained_moedl,
5     environment = e_trade_gym)

```

hit end!

```

1
2 trained_moedl = trained_ddpg
3 df_account_value_ddpg, df_actions_ddpg = DRLAgent.DRL_prediction(
4     model=trained_moedl,
5     environment = e_trade_gym)

```

hit end!

```

1 trained_moedl = trained_ppo
2 df_account_value_ppo, df_actions_ppo = DRLAgent.DRL_prediction(
3     model=trained_moedl,
4     environment = e_trade_gym)

```

hit end!

```

1 trained_moedl = trained_td3
2 df_account_value_td3, df_actions_td3 = DRLAgent.DRL_prediction(
3     model=trained_moedl,
4     environment = e_trade_gym)

```

hit end!

```

1 trained_moedl = trained_sac
2 df_account_value_sac, df_actions_sac = DRLAgent.DRL_prediction(
3     model=trained_moedl,
4     environment = e_trade_gym)

```

hit end!

```
1 df_account_value_a2c.shape
```

(353, 2)

## Part 6.5: Mean Variance Optimization

```
1 mvo_df.head()
```

	date	tic	close
0	2010-01-04	AAPL	6.461976
1	2010-01-04	AMGN	41.200794
2	2010-01-04	AXP	33.090431
3	2010-01-04	BA	43.777538
4	2010-01-04	CAT	40.190208

```
1 fst = mvo_df
2 fst = fst.iloc[0*29:0*29+29, :]
3 tic = fst['tic'].tolist()
4
5 mvo = pd.DataFrame()
6
7 for k in range(len(tic)):
8     mvo[tic[k]] = 0
9
10 for i in range(mvo_df.shape[0]//29):
11     n = mvo_df
12     n = n.iloc[i*29:i*29+29, :]
13     date = n['date'][i*29]
14     mvo.loc[date] = n['close'].tolist()
```

```
1 mvo.shape[0]
```

```
→ 3310
```

## ⌄ Helper functions

Double-click (or enter) to edit

```

1 from scipy import optimize
2 from scipy.optimize import linprog
3
4 #function obtains maximal return portfolio using linear programming
5
6 def MaximizeReturns(MeanReturns, PortfolioSize):
7
8     #dependencies
9
10
11    c = (np.multiply(-1, MeanReturns))
12    A = np.ones([PortfolioSize,1]).T
13    b=[1]
14    res = linprog(c, A_ub = A, b_ub = b, bounds = (0,1), method = 'simplex')
15
16    return res
17
18 def MinimizeRisk(CovarReturns, PortfolioSize):
19
20    def f(x, CovarReturns):
21        func = np.matmul(np.matmul(x, CovarReturns), x.T)
22        return func
23
24    def constraintEq(x):
25        A=np.ones(x.shape)
26        b=1
27        constraintVal = np.matmul(A,x.T)-b
28        return constraintVal
29
30    xinit=np.repeat(0.1, PortfolioSize)
31    cons = ({'type': 'eq', 'fun':constraintEq})
32    lb = 0
33    ub = 1
34    bnds = tuple([(lb,ub) for x in xinit])
35
36    opt = optimize.minimize (f, x0 = xinit, args = (CovarReturns), bounds = bnds, \
37                           constraints = cons, tol = 10**-3)
38
39    return opt
40
41 def MinimizeRiskConstr(MeanReturns, CovarReturns, PortfolioSize, R):
42
43    def f(x,CovarReturns):
44
45        func = np.matmul(np.matmul(x,CovarReturns ), x.T)
46        return func
47
48    def constraintEq(x):
49        AEq=np.ones(x.shape)
50        bEq=1
51        EqconstraintVal = np.matmul(AEq,x.T)-bEq
52        return EqconstraintVal
53
54    def constraintIneq(x, MeanReturns, R):
55        AIneq = np.array(MeanReturns)
56        bIneq = R
57        IneqconstraintVal = np.matmul(AIneq,x.T) - bIneq
58        return IneqconstraintVal
59
60
61    xinit=np.repeat(0.1, PortfolioSize)
62    cons = ({'type': 'eq', 'fun':constraintEq},
63             {'type':'ineq', 'fun':constraintIneq, 'args':(MeanReturns,R) })
64    lb = 0
65    ub = 1
66    bnds = tuple([(lb,ub) for x in xinit])
67
68    opt = optimize.minimize (f, args = (CovarReturns), method ='trust-constr', \
69                           x0 = xinit, bounds = bnds, constraints = cons, tol = 10**-3)
70
71    return opt

```

```

1 def StockReturnsComputing(StockPrice, Rows, Columns):
2     import numpy as np
3     StockReturn = np.zeros([Rows-1, Columns])
4     for j in range(Columns):          # j: Assets
5         for i in range(Rows-1):      # i: Daily Prices
6             StockReturn[i,j]=((StockPrice[i+1, j]-StockPrice[i,j])/StockPrice[i,j])* 100
7
8     return StockReturn

```

## ▼ Calculate mean returns and variance-covariance matrix

```

1 # Obtain optimal portfolio sets that maximize return and minimize risk
2
3 #Dependencies
4 import numpy as np
5 import pandas as pd
6
7
8 #input k-portfolio 1 dataset comprising 15 stocks
9 # StockFileName = './DJIA_Apr112014_Apr112019_kpf1.csv'
10
11 Rows = 1259 #excluding header
12 Columns = 15 #excluding date
13 portfolioSize = 29 #set portfolio size
14
15 #read stock prices in a dataframe
16 # df = pd.read_csv(StockFileName, nrows= Rows)
17
18 #extract asset labels
19 # assetLabels = df.columns[1:Columns+1].tolist()
20 # print(assetLabels)
21
22 #extract asset prices
23 # StockData = df.iloc[0:, 1:]
24 StockData = mvo.head(mvo.shape[0]-336)
25 TradeData = mvo.tail(336)
26 # df.head()
27 TradeData.to_numpy()

```

```

array([[147.03074646, 191.0255127 , 174.92080688, ..., 45.36961746,
       41.60421753, 47.56257629],
       [146.56797791, 189.10319519, 172.09814453, ..., 44.98498917,
       40.57939911, 47.17249298],
       [150.23091125, 189.64329529, 168.79208374, ..., 45.13029099,
       40.56217575, 47.46665573],
       ...,
       [148.40548706, 228.26695251, 172.09477234, ..., 35.28199005,
       33.37779236, 46.33304214],
       [145.73339844, 224.46282959, 171.22024536, ..., 35.47428513,
       32.99985886, 46.45695114],
       [146.93536377, 225.22172546, 170.28675842, ..., 35.60248566,
       32.62192917, 46.12108994]])

```

```

1 #compute asset returns
2 arStockPrices = np.asarray(StockData)
3 [Rows, Cols]=arStockPrices.shape
4 arReturns = StockReturnsComputing(arStockPrices, Rows, Cols)
5
6
7 #compute mean returns and variance covariance matrix of returns
8 meanReturns = np.mean(arReturns, axis = 0)
9 covReturns = np.cov(arReturns, rowvar=False)
10
11 #set precision for printing results
12 np.set_printoptions(precision=3, suppress = True)
13
14 #display mean returns and variance-covariance matrix of returns
15 print('Mean returns of assets in k-portfolio 1\n', meanReturns)
16 print('Variance-Covariance matrix of returns\n', covReturns)

```

```

1.209 1.36  0.821 1.822 0.91  0.925 1.413 0.898 1.249 1.301 0.694 1.25
1.255 1.407 0.694 1.142 0.575]
[1.03  0.851 1.343 1.532 1.355 1.185 1.245 1.275 1.058 1.317 0.99  1.209
2.009 1.281 0.706 1.37  0.77  0.749 1.08  0.782 1.123 0.947 0.664 0.995
1.003 1.106 0.644 0.958 0.541]
[1.525 1.153 1.508 1.699 1.566 1.637 1.583 1.386 1.228 1.564 1.243 1.36
1.281 3.282 0.787 1.563 0.786 0.798 1.18  0.857 1.695 1.123 0.722 1.122
1.148 1.329 0.674 1.083 0.7 ]
[0.692 0.891 0.861 0.905 0.826 0.756 0.843 0.836 0.701 0.83  0.691 0.827
0.706 0.787 1.121 0.878 0.611 0.561 0.788 0.777 0.784 0.683 0.637 0.693
0.833 0.795 0.566 0.779 0.521]
[1.269 1.14  2.296 2.039 1.564 1.511 1.942 1.68  2.724 1.364 1.822
1.37  1.563 0.878 3.233 0.93  0.961 1.43  1.024 1.376 1.344 0.737 1.621
1.4  1.559 0.807 1.285 0.635]
[0.664 0.641 0.99  1.117 0.848 0.738 0.728 0.919 0.829 0.825 0.73  0.91
0.77  0.786 0.611 0.93  1.195 0.653 0.759 0.634 0.744 0.733 0.689 0.832
0.735 0.807 0.59  0.681 0.489]
[0.816 0.664 1.075 1.195 0.855 0.984 0.823 0.993 0.863 0.939 0.91  0.925
0.749 0.798 0.561 0.961 0.653 1.421 0.75  0.641 0.874 0.907 0.543 0.888
0.808 0.943 0.481 0.605 0.468]
[1.029 0.906 1.375 1.461 1.585 1.151 1.215 1.263 1.117 1.41  1.032 1.413
1.08  1.18  0.788 1.43  0.759 0.75  1.822 0.755 1.027 1.036 0.649 1.03
1.004 1.119 0.63  1.024 0.557]
[0.733 1.072 0.964 0.949 0.918 0.91  0.881 0.963 0.775 0.959 0.762 0.898
0.782 0.857 0.777 1.024 0.634 0.641 0.755 1.69  0.853 0.754 0.621 0.821
0.921 0.888 0.629 0.784 0.518]
[1.576 1.088 1.357 1.437 1.331 1.853 1.481 1.203 1.15  1.418 1.205 1.249
1.123 1.695 0.784 1.376 0.744 0.874 1.027 0.853 2.495 1.16  0.751 0.989
1.19  1.407 0.636 0.963 0.704]
[1.139 0.873 1.418 1.573 1.292 1.493 1.145 1.136 1.197 1.32  1.202 1.301
0.947 1.123 0.683 1.344 0.733 0.907 1.036 0.754 1.16  2.688 0.628 0.981
1.086 1.275 0.561 0.879 0.582]
[0.668 0.724 0.728 0.748 0.682 0.657 0.74  0.686 0.663 0.693 0.691 0.694
0.664 0.722 0.637 0.737 0.689 0.543 0.649 0.621 0.751 0.628 1.146 0.676
0.666 0.69  0.576 0.689 0.601]
[0.86  0.854 1.505 1.604 1.269 1.042 1.002 1.361 1.097 1.496 1.078 1.25
0.995 1.122 0.693 1.621 0.832 0.888 1.03  0.821 0.989 0.981 0.676 1.974
1.089 1.116 0.7  0.922 0.588]
[1.144 1.146 1.417 1.472 1.221 1.204 1.098 1.296 1.058 1.383 1.14  1.255
1.003 1.148 0.833 1.4  0.735 0.808 1.004 0.921 1.19  1.086 0.666 1.089
2.628 1.185 0.61  1.003 0.618]
[1.338 1.056 1.8  1.716 1.408 1.722 1.31  1.381 1.313 1.53  1.207 1.407
1.106 1.329 0.795 1.559 0.807 0.943 1.119 0.888 1.407 1.275 0.69  1.116
1.185 2.579 0.606 0.924 0.572]
[0.525 0.631 0.748 0.734 0.739 0.563 0.655 0.74  0.66  0.717 0.628 0.694
0.644 0.674 0.566 0.807 0.59  0.481 0.63  0.629 0.636 0.561 0.576 0.7
0.61  0.606 1.194 0.662 0.51 ]
[0.908 1.025 1.216 1.362 1.226 1.017 1.065 1.06  0.994 1.231 0.915 1.142
0.958 1.083 0.779 1.285 0.681 0.605 1.024 0.784 0.963 0.879 0.689 0.922
1.003 0.924 0.662 3.14  0.709]
[0.635 0.645 0.584 0.603 0.613 0.621 0.686 0.518 0.574 0.615 0.723 0.575
0.541 0.7  0.521 0.635 0.489 0.468 0.557 0.518 0.704 0.582 0.601 0.588
0.618 0.572 0.51  0.709 1.406]]

```

```

1 from pypfopt.efficient_frontier import EfficientFrontier
2
3 ef_mean = EfficientFrontier(meanReturns, covReturns, weight_bounds=(0, 0.5))
4 raw_weights_mean = ef_mean.max_sharpe()
5 cleaned_weights_mean = ef_mean.clean_weights()
6 mvo_weights = np.array([100000 * cleaned_weights_mean[i] for i in range(29)])
7 mvo_weights

```

```

array([268510.,      0.,      0.,      0.,   67170.,      0.,
       0.,      0., 328490.,      0.,      0.,      0.,
       0.,      0.,      0.,      0.,      0.,      0.,
     80330.,      0.,      0., 249920.,   5570.,      0.,      0.,
       0.])

```

```
1 StockData.tail(1)
```

	AAPL	AMGN	AXP	BA	CAT	CRM	CSCO	CVX	DIS	GS	...	MRK
2021-10-25	146.36116	190.064377	176.235519	212.869995	192.009811	293.53656	51.05426	102.366913	171.451218	384.550415	...	75.308884

1 rows × 29 columns

```

1 LastPrice = np.array([1/p for p in StockData.tail(1).to_numpy()[0]])
2 Initial_Portfolio = np.multiply(mvo_weights, LastPrice)
3 Initial_Portfolio

→ array([1834.571,    0.    ,    0.    ,    0.    ,    0.    ,  228.83 ,
       0.    ,    0.    ,    0.    ,    0.    ,  948.3  ,    0.    ,
       0.    ,    0.    ,    0.    ,    0.    ,    0.    ,    0.    ,
       0.    ,    0.    ,    0.    ,  505.467,    0.    ,    0.    ,
      578.174,  24.319,    0.    ,    0.    ,    0.    ])

```

```

1 Portfolio_Assets = TradeData @ Initial_Portfolio
2 MVO_result = pd.DataFrame(Portfolio_Assets, columns=["Mean Var"])
3 MVO_result

```

	Mean Var
2021-10-26	1.002272e+06
2021-10-27	1.002461e+06
2021-10-28	1.011236e+06
2021-10-29	1.010909e+06
2021-11-01	1.001450e+06
...	...
2023-02-21	9.195315e+05
2023-02-22	9.194654e+05
2023-02-23	9.249119e+05
2023-02-24	9.120102e+05
2023-02-27	9.135294e+05

336 rows × 1 columns

+ Code + Text

## Part 7: Backtesting Results

```

1 df_result_a2c = df_account_value_a2c.set_index(df_account_value_a2c.columns[0])
2 df_result_ddpg = df_account_value_ddpg.set_index(df_account_value_ddpg.columns[0])
3 df_result_td3 = df_account_value_td3.set_index(df_account_value_td3.columns[0])
4 df_result_ppo = df_account_value_ppo.set_index(df_account_value_ppo.columns[0])
5 df_result_sac = df_account_value_sac.set_index(df_account_value_sac.columns[0])
6 df_account_value_a2c.to_csv("df_account_value_a2c.csv")
7 #baseline stats
8 print("=====Get Baseline Stats====")
9 df_dji_ = get_baseline(
10     ticker="^DJI",
11     start = TRADE_START_DATE,
12     end = TRADE_END_DATE)
13 stats = backtest_stats(df_dji_, value_col_name = 'close')
14 df_dji = pd.DataFrame()
15 df_dji['date'] = df_account_value_a2c['date']
16 df_dji['account_value'] = df_dji_[['close']] / df_dji_[['close']][0] * env_kwargs["initial_amount"]
17 df_dji.to_csv("df_dji.csv")
18 df_dji = df_dji.set_index(df_dji.columns[0])
19 df_dji.to_csv("df_dji+.csv")
20
21 result = pd.merge(df_result_a2c, df_result_ddpg, left_index=True, right_index=True, suffixes=('_a2c', '_ddpg'))
22 result = pd.merge(result, df_result_td3, left_index=True, right_index=True, suffixes=('', '_td3'))
23 result = pd.merge(result, df_result_ppo, left_index=True, right_index=True, suffixes=('', '_ppo'))
24 result = pd.merge(result, df_result_sac, left_index=True, right_index=True, suffixes=('', '_sac'))
25 result = pd.merge(result, MVO_result, left_index=True, right_index=True, suffixes=('', '_mvo'))
26 result = pd.merge(result, df_dji, left_index=True, right_index=True, suffixes=('', '_dji'))
27 result.columns = ['a2c', 'ddpg', 'td3', 'ppo', 'sac', 'mean var', 'dji']
28
29 print("result: ", result)
30 result.to_csv("result.csv")

→ =====Get Baseline Stats=====
[*****100%*****] 1 of 1 completedShape of DataFrame: (354, 8)
Annual return      -0.034876
Cumulative returns -0.048644
Annual volatility   0.181612

```

```

Sharpe ratio      -0.105351
Calmar ratio     -0.158953
Stability        0.280983
Max drawdown    -0.219408
Omega ratio       0.982546
Sortino ratio    -0.146974
Skew              NaN
Kurtosis          NaN
Tail ratio        0.970602
Daily value at risk -0.022957
dtype: float64

result:           a2c      ddpg      td3      ppo   \
2021-10-26  1.025094e+06  1.052065e+06  1.025573e+06  1.019773e+06
2021-10-27  1.013737e+06  1.050434e+06  1.006178e+06  1.013452e+06
2021-10-28  1.023732e+06  1.060355e+06  1.014938e+06  1.018103e+06
2021-10-29  1.018214e+06  1.062797e+06  1.013954e+06  1.017476e+06
2021-11-01  1.022763e+06  1.064088e+06  1.019041e+06  1.020429e+06
...
...
...
...
2023-02-21  1.039334e+06  1.041966e+06  1.031807e+06  1.011313e+06
2023-02-22  1.037995e+06  1.038825e+06  1.027912e+06  1.011083e+06
2023-02-23  1.040820e+06  1.042367e+06  1.031837e+06  1.013820e+06
2023-02-24  1.029576e+06  1.034298e+06  1.024170e+06  1.006364e+06
2023-02-27  1.034299e+06  1.037987e+06  1.028619e+06  1.008611e+06

sac      mean var      dji
2021-10-26  1.038037e+06  1.002272e+06  1.041671e+06
2021-10-27  1.038085e+06  1.002461e+06  1.033916e+06
2021-10-28  1.043188e+06  1.011236e+06  1.040902e+06
2021-10-29  1.046913e+06  1.010909e+06  1.043497e+06
2021-11-01  1.052176e+06  1.001450e+06  1.045244e+06

```